

Jeremy Goldwasser

Website · Email · LinkedIn · Google Scholar

EDUCATION	UC Berkeley , Berkeley, CA	Aug 2021 - May 2026 (Expected)
	<i>Ph.D. in Statistics; Advised by Ryan Tibshirani and Giles Hooker</i>	GPA: 3.97
	Yale University , New Haven, CT	Aug 2017 - May 2021
	<i>B.S. in Statistics and Data Science</i>	GPA: 3.73
INTERESTS	Interpretability; statistical ML; uncertainty quantification; time series; LLMs; AI for health.	
INTERNSHIP EXPERIENCE	Apple	May - Aug 2024
	<i>Research Intern, Vision Products Group</i>	Sunnyvale, CA
	- Developed novel generative AI method for counterfactual image explanations. - Tool is currently used to provide previously unattainable insights on failure modes of FaceID.	
	Genentech	May - July 2023
	<i>Data and Statistical Sciences Intern</i>	South San Francisco, CA
	- Produced confidence intervals for AI pathology models using conformal prediction. - Calibrated cell-type classifier, quantifying uncertainty and improving test accuracy by 3%.	
	Voca.ai (acquired by Snapchat)	June - Aug 2019
	<i>Machine Learning Intern</i>	Herzliya, Israel
	Trained Transformers for Automatic Speech Recognition and Named Entity Recognition.	
PUBLICATIONS	Unifying Image Counterfactuals and Feature Attributions with Latent-Space Adversarial Attacks	
	<i>ICML 2025, Actionable Interpretability Workshop</i> – First author	
	Statistical Significance of Feature Importance Rankings	
	<i>Conference on Uncertainty in AI (UAI) 2025</i> – First author	
	Gaussian Rank Verification	
	<i>Stat, 2025</i> – First author	
	Stabilizing Estimates of Shapley Values with Control Variates	
	<i>World Conference on Explainability in AI (XAI) 2024</i> – First author	
	Ascle: A Python Natural Language Processing Toolkit for Medical Text Generation	
	<i>Journal of American Medical Informatics Association (JAMIA), 2024</i>	
	Forest Fire Clustering for Single-Cell Sequencing Combines Iterative Label Propagation with Parallelized Monte Carlo Simulations	
	<i>Nature Communications, 2022</i> – Second author	
	Neural NLP for Unstructured Data in Electronic Health Records: A Review	
	<i>Computer Science Review, 2022</i>	
	Challenges in Real-Time Estimation of Changing Epidemic Severity Rates	
	<i>In Review, PLOS Computational Biology</i> – First author, 2024	
ONGOING PROJECTS	Developing time series algorithms to track mortality risk in epidemics (first author). Our statistical ML approach improves accuracy by 40–65% over widely used baselines.	
	- Attention-based interpretability of Vision Transformers for AI cancer pathology. - Online learning for ensembling time series forecasters with uncertainty quantification.	

SKILLS	Coursework: Deep Learning, Optimization, Causal Inference, AI in Biology Programming: Python (PyTorch, NumPy, scikit-learn, Pandas, Matplotlib), R, L ^A T _E X Languages: Spanish (fluent), Hebrew (basic)	
MISC PROJECTS	AI Meets DNA Methylation May 2024 - Used interpretable ML to identify methylation sites that regulate CD55 gene expression. Predicting Peptide-MHC Binding Affinity in SARS-CoV-2 April 2020 - Designed neural architecture to predict which viral peptides bind with immune cells.	
AWARDS	U.S. Civic Digital Fellowship (Declined) 2021 Yale College Dean's Research Fellowship, Morse Richter Fellowship 2020 Yale Creative and Performing Arts Grant 2020	
INVITED TALKS	“Estimating Epidemic Severity Rates.” UCSF MINDSCAPE, February 2025. “Estimating Epidemic Severity Rates.” Delphi Group, January 2025. “Estimation of Time-Varying Severity Rates from Aggregate Data: Pitfalls and New Ideas.” California Department of Public Health, September 2024. “Estimation of Time-Varying Severity Rates from Aggregate Data: Pitfalls and New Ideas.” Centers for Disease Control and Prevention (CDC), September 2024. “Stabilizing Estimates of Shapley Values with Control Variates.” 2nd World Conference on Explainable AI, Malta, July 2024. “Provably Stable Feature Rankings with SHAP and LIME.” Apple Data Analytics and Quality Group, June 2024. “Stable Feature Rankings & Time Series Deconvolution.” Genentech gRED Computational Sciences, February 2024. “Provably Stable Feature Rankings with SHAP and LIME.” Center for Human-Compatible AI, January 2024. “Provably Stable Feature Rankings with SHAP and LIME.” Berkeley Statistics Student Seminar, January 2024. “Estimating Case Fatality Rate via Convolutional Modeling.” Delphi Group, April 2023.	